



Zachary Huang

Posted on Mar 31, 2025



# Retrieval Augmented Generation (RAG) from Scratch — Tutorial For Dummies

#ai #programming #python #rag

*Ever wondered how AI tools like ChatGPT can answer questions based on specific documents they've never seen before? This guide breaks down Retrieval Augmented Generation (RAG) in the simplest possible way with [minimal code implementation!](#)*

Have you ever asked an AI a question about your personal documents and received a completely made-up answer? Or maybe you've been frustrated when chatbots provide outdated information? These are common problems with traditional LLMs (Large Language Models), but there's a solution: **Retrieval Augmented Generation (RAG)**.

In this beginner-friendly tutorial, you'll learn:

- The key concepts behind RAG systems in plain language
- How to build a working RAG system step-by-step
- Why RAG dramatically improves AI responses for your specific documents

We'll use [\*\*PocketFlow\*\*](#) - a simple 100-line framework that strips away complexity. Unlike other frameworks with convoluted abstractions, PocketFlow lets you see the entire system at once, giving you the fundamentals to build your understanding from the ground up.

## What's RAG (In Human Terms)?

Imagine RAG is like giving an AI its own personal research librarian before it answers your questions. Here's how the magic happens:

1. **Document Collection:** You provide your documents (company manuals, articles, books) to the system, just like books being added to a library.
2. **Chunking:** The system breaks these down into bite-sized, digestible pieces - like librarians dividing books into chapters and sections rather than working with entire volumes.
3. **Embedding:** Each chunk gets converted into a special numerical format (vector) that captures its meaning - similar to creating detailed index cards that understand concepts, not just keywords.
4. **Indexing:** These vectors are organized in a searchable database - like a magical card catalog that understands the relationships between different topics.
5. **Retrieval:** When you ask a question, the system consults its index to find the most relevant chunks related to your query.
6. **Generation:** The AI crafts an answer using both your question AND these helpful references, producing a much better response than if it relied solely on its pre-trained knowledge.

The result? Instead of making things up or giving outdated information, the AI grounds its answers in your specific documents, providing accurate, relevant responses tailored to your information.

## Chunking: Breaking Documents into Manageable Pieces

Before our RAG system can work effectively, we need to break our documents into smaller, digestible pieces. Think of chunking like preparing a meal for guests - you wouldn't serve a whole turkey without carving it first!

## Why Chunking Matters

The size of your chunks directly impacts the quality of your RAG system:

- **Too large chunks:** Your system retrieves too much irrelevant information (like serving entire turkeys)
- **Too small chunks:** You lose important context (like serving single peas)
- **Just right chunks:** Your system finds precisely what it needs (perfect portions!)

Let's explore some practical chunking methods:

### 1. Fixed-Size Chunking: Simple but Imperfect

The simplest approach divides text into equal-sized pieces, regardless of content:

```
def fixed_size_chunk(text, chunk_size=50):
    chunks = []
    for i in range(0, len(text), chunk_size):
        chunks.append(text[i:i+chunk_size])
    return chunks
```

This code loops through text, taking 50 characters at a time. Let's see how it works on a sample paragraph:

#### Input Text:

The quick brown fox jumps over the lazy dog. Artificial intelligence has revolutionized many industries. Today's weather is sunny with a chance of rain. Many researchers work on RAG systems to improve information retrieval.

#### Output Chunks:

Chunk 1: "The quick brown fox jumps over the lazy dog. Arti"  
 Chunk 2: "ficial intelligence has revolutionized many indus"  
 Chunk 3: "tries. Today's weather is sunny with a chance of "  
 Chunk 4: "rain. Many researchers work on RAG systems to imp"  
 Chunk 5: "rove information retrieval."

Notice the problem? The word "Artificial" is split between chunks 1 and 2.

"Industries" is split between chunks 2 and 3. This makes it hard for our system to understand the content properly.

### 2. Sentence-Based Chunking: Respecting Natural Boundaries

A smarter approach is to chunk by complete sentences:

```
import nltk # Natural Language Toolkit library

def sentence_based_chunk(text, max_sentences=1)
    sentences = nltk.sent_tokenize(text)
    chunks = []
    # Group sentences, 1 at a time
    for i in range(0, len(sentences), max_sentences):
        chunks.append(" ".join(sentences[i:i+max_sentences]))
    return chunks
```

This method first identifies complete sentences, then groups them one at a time:

### Output Chunks:

```
Chunk 1: "The quick brown fox jumps over the lazy dog."
Chunk 2: "Artificial intelligence has revolutionized many industries."
Chunk 3: "Today's weather is sunny with a chance of rain."
Chunk 4: "Many researchers work on RAG systems to improve information retrieval"
```

Much better! Each chunk now contains a complete sentence with its full meaning intact.

## 3. Additional Chunking Strategies

Depending on your documents, these approaches might also work well:

- **Paragraph-Based:** Split text at paragraph breaks (usually marked by newlines)
- **Semantic Chunking:** Group text by topics or meaning (often requires AI assistance)
- **Hybrid Approaches:** Combine multiple strategies for optimal results

## Choosing the Right Approach

While many sophisticated chunking methods exist, for most practical applications, the "Keep It Simple, Stupid" (KISS) principle applies. **Starting with Fixed-Size Chunking of around 1,000 characters per chunk is often sufficient** and avoids overcomplicating your system.

The best chunking approach ultimately depends on your specific documents and use case - just like a chef adjusts portions based on the meal and guests!

## Embeddings: Making Retrieval Possible

Now that we've chopped up our documents, how does the system find the most relevant chunks for our questions? This is where **embeddings** come in! Embeddings are the magic that powers our retrieval system. Without them, we couldn't find the right information for our questions!

## What Are Embeddings?

An embedding transforms text into a list of numbers (a vector) that captures its meaning. Think of it as creating a special "location" in a meaning-space where similar ideas are positioned close together.

For example, if we take these three sentences:

1. "The cat sat on the mat." 🐱
2. "A feline rested on the floor covering." 🐾
3. "Python is a popular programming language." 💻

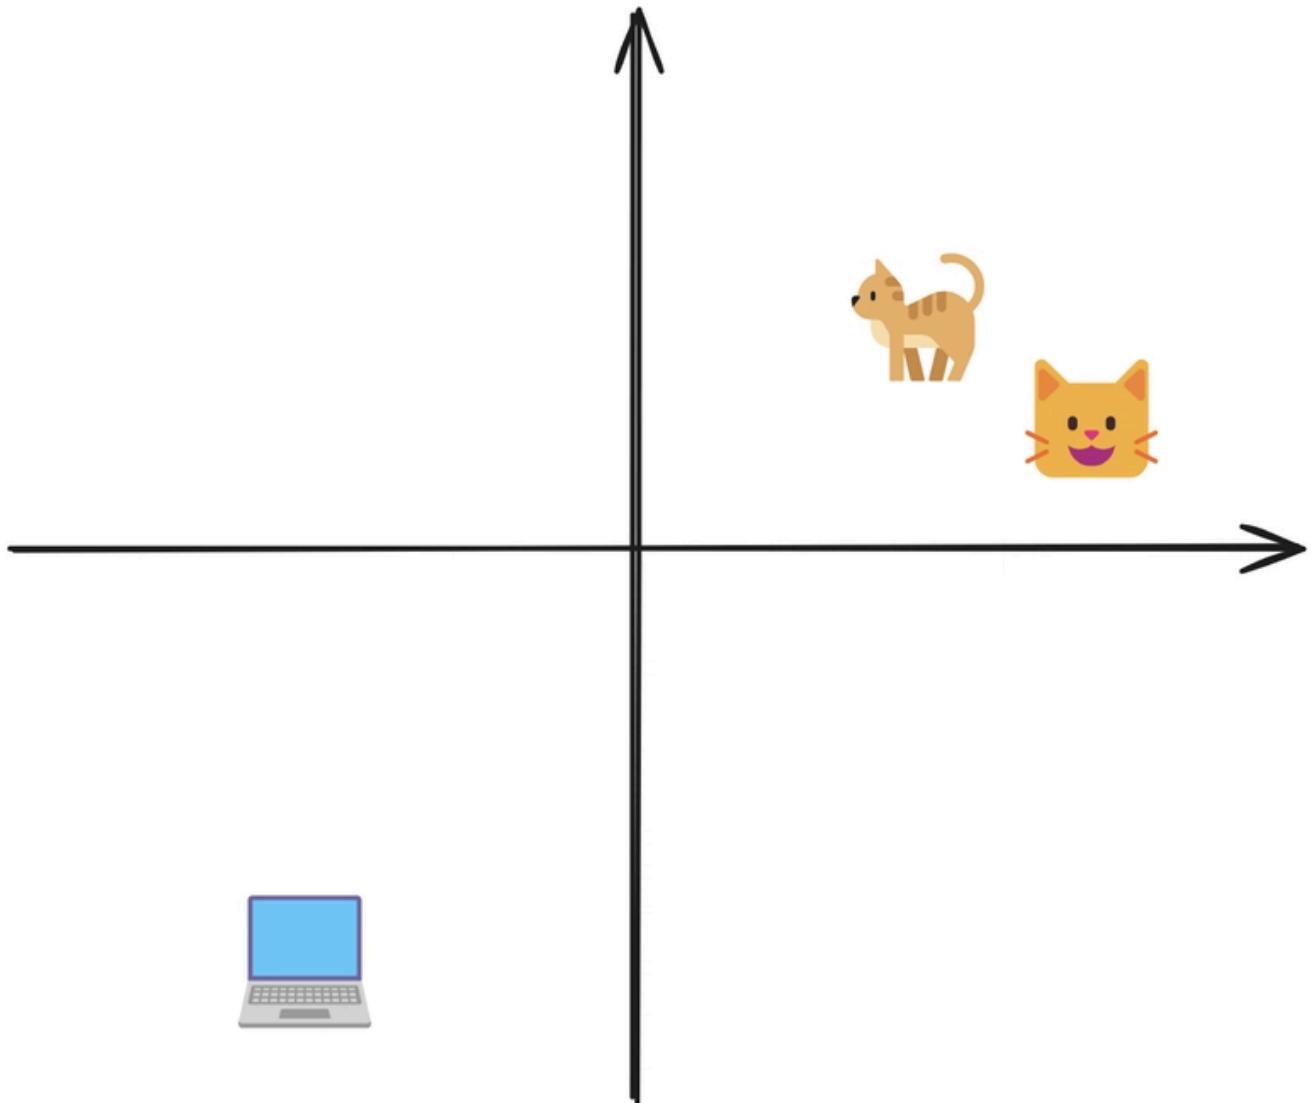
A good embedding system would place sentences 1 and 2 close together (since they describe the same concept), while sentence 3 would be far away (completely different topic).

## Visualizing Embeddings in Action

When we convert these sentences to embeddings, we might get something like this in a simplified 2D space:

```
Sentence 1 (🐱): [0.8, 0.2]
Sentence 2 (🐯): [0.7, 0.3]
Sentence 3 (💻): [-0.5, -0.9]
```

On a graph, the cat-related sentences would cluster in the upper right, while the programming sentence would sit in the lower left - visually showing their semantic relationships!



With embeddings, measuring similarity is as simple as calculating the distance between points:

- Distance from 🐱 to 🐈: Very small (about 0.14 units) → Very similar!
- Distance from 🐱 to 💻: Very large (about 1.95 units) → Not related at all!

This confirms what we intuitively know – the cat sentences are closely related, while the programming sentence is completely different.

## Creating Embeddings: From Simple to Advanced

### Simple Approach: Character Frequency

Here's a beginner-friendly way to create embeddings by counting characters:

```
def get_simple_embedding(text):  
    # Create a vector to store character frequencies  
    embedding = np.zeros(26, dtype=np.float32)  
  
    # Count character frequencies in the text  
    for char in text:
```

```
embedding[ord(char) % 26] += 1.0
```

```
# Normalize the vector
norm = np.linalg.norm(embedding)
if norm > 0:
    embedding /= norm

return embedding
```

Let's see what this looks like for our three example sentences:

 "The cat sat on the mat." → [0.2, 0, 0, ...]  
(high values at positions for 't', 'a', ' ', etc.)

 "A feline rested on the floor covering." → [0.1, 0.2, 0, ...]  
(different characters but similar number of spaces)

 "Python is a popular programming language." → [0.1, 0.2, 0.1, ...]  
(completely different character pattern)

Notice how sentences 1 and 2 have different values because they use different words, while sentence 3 has a very different pattern altogether. This simple vector only captures the characters used, not their meaning.

**Limitation:** This simple approach can't recognize that "cat" and "feline" are related concepts since they share no characters!

## Professional Approach: AI-Powered Embeddings

In a real RAG system, you'd use sophisticated models like OpenAI's embedding API:

```
def get_openai_embedding(text):
    client = OpenAI(api_key=os.environ.get("OPENAI_API_KEY", "YOUR_API_KEY"))

    response = client.embeddings.create(
        model="text-embedding-ada-002",
        input=text
    )

    # Extract the embedding vector from the response
    embedding = response.data[0].embedding

    return np.array(embedding, dtype=np.float32)
```

These advanced embeddings capture deep semantic relationships in high-dimensional space (typically 1,536 dimensions for OpenAI's embeddings). They understand that "cat" 🐱 and "feline" 🐾 mean the same thing, even with completely different characters!

🐱 "The cat sat on the mat." → [0.213, -0.017, 0.122, ...]  
(pattern encoding "animal resting on object")

🐯 "A feline rested on the floor covering." → [0.207, -0.024, 0.118, ...]  
(very similar pattern, also encoding "animal resting on object")

💻 "Python is a popular programming language." → [-0.412, 0.158, 0.367, ...]  
(completely different pattern encoding "programming language")

Even though sentences 1 and 2 use different words, their embeddings have similar patterns because they express similar concepts. Meanwhile, sentence 3's embedding has a completely different pattern because it's about a different topic entirely.

## How Embeddings Drive Our RAG System

In our RAG pipeline:

1. We embed all document chunks and store these vectors
2. When a user asks a question, we embed the question too
3. We find the document chunks whose embeddings are closest to the question's embedding

However, as your document collection grows into thousands or millions of chunks, searching through all embeddings becomes slow. That's where vector databases come in - specialized systems designed to make this search lightning-fast!

## Vector Databases: Making Retrieval Fast

Imagine having to search through a million book pages to find an answer - it would take forever! Vector databases solve this problem by organizing our embeddings in a clever way that makes searching lightning-fast.

## Why We Need Vector Databases

When you ask a question, your RAG system needs to compare it against potentially thousands or millions of document chunks. There are two ways to do this:

## 1. Simple Approach: The Exhaustive Search

In this approach, we check every single document chunk one by one - like going through every page in a library:

```
def retrieve_naive(question_embedding, chunk_embeddings):
    best_similarity, best_chunk_index = -1, -1

    for idx, chunk_embedding in enumerate(chunk_embeddings):
        similarity = get_similarity(question_embedding, chunk_embedding)
        if similarity > best_similarity:
            best_similarity, best_chunk_index = similarity, idx

    return best_chunk_index
```

This works perfectly for a few dozen documents, but becomes painfully slow with thousands or millions of chunks.

## 2. Professional Approach: Smart Indexing with Vector Databases

Vector databases are like having a magical librarian who organizes books so effectively that they can instantly find what you need without checking every shelf.

Let's see how this works in two simple steps:

### Step 1: Building the Magical Index

First, we organize all our document embeddings into a special structure:

```
def create_index(chunk_embeddings):
    dimension = chunk_embeddings.shape[1] # e.g. 128 or 1536
    index = faiss.IndexFlatL2(dimension) # flat = exact search
    index.add(chunk_embeddings) # add all document vectors
    return index
```

This is like creating a detailed map of where every document "lives" in our meaning-space.

### Step 2: Fast Retrieval Using the Index

When a question comes in, we use our magical index to find the most relevant chunks in an instant:

```
def retrieve_index(index, question_embedding, top_k=5):
    _, chunk_indices = index.search(question_embedding, top_k)
    return chunk_indices
```

Instead of checking every document, the index knows exactly where to look - giving us the top 5 most relevant chunks in milliseconds!

## What Makes Vector Databases So Efficient?

The speed of vector databases comes from three clever techniques:

- 1. Smart Indexing Algorithms:** Methods like **HNSW** (Hierarchical Navigable Small Worlds) create shortcuts through the embedding space, so the system only needs to check a small fraction of documents.
- 2. Vector Compression:** These databases can shrink embeddings to save memory while preserving their relationships - like having a compressed map that still shows all the important landmarks.
- 3. Parallel Processing:** Modern vector databases use multiple CPU/GPU cores simultaneously, checking many possibilities at once - like having a team of librarians all searching different sections of the library at the same time.

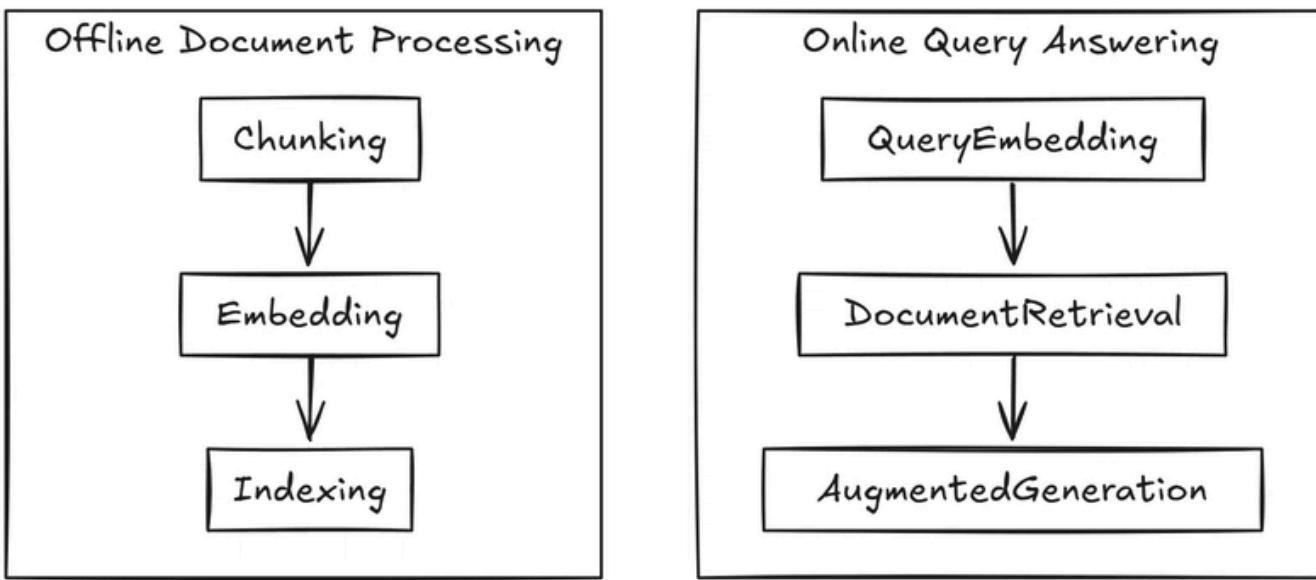
With these techniques, vector databases can search through millions of documents in milliseconds - making RAG systems practical for even the largest document collections!

## Putting RAG Together: Just Two Simple Workflows

Now that you understand chunking, embeddings, and vector databases, here's the beautiful simplicity of RAG that many frameworks overcomplicate:

**RAG is just two straightforward workflows working together!**

Think of RAG like your personal research assistant, organized into two efficient processes:



## 1. Offline Flow: Preparing Your Knowledge Base

This happens just once, before any questions are asked:

- **Chunking:** Break documents into bite-sized pieces (like dividing a book into memorable paragraphs)
- **Embedding:** Convert each chunk into a numerical vector (like giving each paragraph a unique "location" in meaning-space)
- **Indexing:** Organize these vectors for efficient retrieval (like creating a magical map of all your knowledge)

## 2. Online Flow: Answering Questions in Real-Time

This happens each time someone asks a question:

- **Query Embedding:** Transform the question into a vector (finding its "location" in the same meaning-space)
- **Retrieval:** Find the document chunks closest to the question (locate the most relevant knowledge)
- **Augmented Generation:** Create an answer using both the question and retrieved context (craft a response grounded in your specific documents)

That's it! The magic of RAG isn't in complex algorithms - it's in how these two simple workflows work together to provide accurate, relevant answers based on your specific knowledge!

## Building RAG with PocketFlow: A Simple Framework

[PocketFlow](#) makes building RAG systems delightfully simple. Unlike complex frameworks with layers of abstraction, PocketFlow uses minimal building blocks in

[just 100 lines of code](#) to clearly show how everything works.

Think of PocketFlow as a well-organized kitchen where:

- **Nodes** are cooking stations that perform specific tasks:

```
class BaseNode:
    def __init__(self):
        self.params, self.successors = {}, {}

    def add_successor(self, node, action="default"):
        self.successors[action] = node
        return node

    def prep(self, shared): pass
    def exec(self, prep_res): pass
    def post(self, shared, prep_res, exec_res): pass

    def run(self, shared):
        p = self.prep(shared)
        e = self.exec(p)
        return self.post(shared, p, e)
```

- **Flow** is the recipe that coordinates these stations:

```
class Flow(BaseNode):
    def __init__(self, start):
        super().__init__()
        self.start = start

    def get_next_node(self, curr, action):
        return curr.successors.get(action or "default")

    def orch(self, shared, params=None):
        curr, p = copy.copy(self.start), (params or {**self.params})
        while curr:
            curr.set_params(p)
            c = curr.run(shared)
            curr = copy.copy(self.get_next_node(curr, c))

    def run(self, shared):
        pr = self.prep(shared)
        self.orch(shared)
        return self.post(shared, pr, None)
```

- **Shared Store** is a common workspace where all nodes can access ingredients:

```
# Connect nodes together
load_data_node = LoadDataNode()
summarize_node = SummarizeNode()
```

```
load_data_node >> summarize_node

# Create flow
flow = Flow(start=load_data_node)

# Pass data through shared store
shared = {"file_name": "data.txt"}
flow.run(shared)
```

Each **Node** performs three simple operations:

- **Prep:** Gather what's needed from the shared store
- **Exec:** Perform its specific task
- **Post:** Store results and decide what happens next

The **Flow** manages the entire process, moving data smoothly from one node to the next based on specific conditions.

With these simple building blocks, PocketFlow makes RAG systems easy to build, understand, and modify!

## Building a Practical RAG System with PocketFlow

Let's create an AI assistant that can answer questions from your documents. Our goal is to build a simple but powerful RAG system that:

1. Processes your documents only once (offline flow)
2. Takes user questions
3. Retrieves the most relevant information
4. Generates accurate, document-grounded answers

## Step-By-Step Walkthrough with a Real Example

Let's see what happens when you ask our system: "**What's the capital of France?**"

### Offline Phase (Happens Once)

#### 1. ChunkDocuments Node:

- **Input ( Prep ):** Your collection of documents
- **Process ( Exec ):** Divides each document into manageable chunks
- **Output ( Post ):** Collection of document chunks (including one with "Paris is the capital of France")

#### 2. EmbedDocuments Node:

- **Input ( Prep )**: Document chunks
- **Process ( Exec )**: Converts each chunk into a numerical vector
- **Output ( Post )**: Collection of document embeddings

### 3. CreateIndex Node:

- **Input ( Prep )**: Document embeddings
- **Process ( Exec )**: Builds a searchable vector index
- **Output ( Post )**: Ready-to-use vector database

## Online Phase (Happens Every Question)

### 1. EmbedQuery Node:

- **Input ( Prep )**: "What's the capital of France?"
- **Process ( Exec )**: Converts question into an embedding vector
- **Output ( Post )**: Question embedding

### 2. RetrieveDocuments Node:

- **Input ( Prep )**: Question embedding + vector database
- **Process ( Exec )**: Finds the most similar document chunk
- **Output ( Post )**: Retrieves "Paris is the capital of France" chunk

### 3. GenerateAnswer Node:

- **Input ( Prep )**: Original question + retrieved chunk
- **Process ( Exec )**: Crafts response using both inputs
- **Output ( Post )**: "The capital of France is Paris."

Now let's implement this with code!

## Building the Offline Flow

### Node 1: The ChunkDocumentsNode

This node breaks documents into smaller, manageable pieces:

```
class ChunkDocumentsNode(BatchNode):
    def prep(self, shared):
        return shared["texts"]

    def exec(self, text):
        return fixed_size_chunk(text)

    def post(self, shared, prep_res, exec_res_list):
```

```

all_chunks = []
for chunks in exec_res_list:
    all_chunks.extend(chunks)

shared["texts"] = all_chunks

print(f"✓ Created {len(all_chunks)} chunks from {len(prep_res)} documents")
return "default"

```

The prep gets documents from shared storage, exec divides each document into chunks, and post combines all chunks into a single list for the next node.

## Node 2: The EmbedDocumentsNode

This node converts each text chunk into a numerical vector:

```

class EmbedDocumentsNode(BatchNode):
    def prep(self, shared):
        return shared["texts"]

    def exec(self, text):
        return get_embedding(text)

    def post(self, shared, prep_res, exec_res_list):
        embeddings = np.array(exec_res_list, dtype=np.float32)

        shared["embeddings"] = embeddings

        print(f"✓ Created {len(embeddings)} document embeddings")
        return "default"

```

The prep retrieves text chunks, exec converts each chunk into an embedding vector, and post stores all embeddings as a single array.

## Node 3: The CreateIndexNode

This node builds a searchable index for fast retrieval:

```

class CreateIndexNode(Node):
    def prep(self, shared):
        return shared["embeddings"]

    def exec(self, embeddings):
        print("🔍 Creating search index...")

        dimension = embeddings.shape[1]

```

```

index = faiss.IndexFlatL2(dimension)

index.add(embeddings)

return index

def post(self, shared, prep_res, exec_res):
    shared["index"] = exec_res

    print(f"✓ Index created with {exec_res.ntotal} vectors")
    return "default"

```

The `prep` gets embeddings, `exec` creates and populates a vector index, and `post` saves the index for query processing.

## Building the Online Flow

### Node 4: The EmbedQueryNode

This node converts a user question into the same vector format:

```

class EmbedQueryNode(Node):
    def prep(self, shared):
        return shared["query"]

    def exec(self, query):
        print(f"🔍 Embedding query: {query}")

        query_embedding = get_embedding(query)

        return np.array([query_embedding], dtype=np.float32)

    def post(self, shared, prep_res, exec_res):
        shared["query_embedding"] = exec_res

        return "default"

```

The `prep` gets the user's question, `exec` creates an embedding using the same method as for documents, and `post` stores this embedding for retrieval.

### Node 5: The RetrieveDocumentNode

This node finds the most relevant document chunks:

```

class RetrieveDocumentNode(Node):
    def prep(self, shared):

```

```

return shared["query_embedding"], shared["index"], shared["texts"]

def exec(self, inputs):
    print("🔍 Searching for relevant documents...")

    query_embedding, index, texts = inputs

    distances, indices = index.search(query_embedding, k=1)

    best_idx = indices[0][0]
    distance = distances[0][0]
    most_relevant_text = texts[best_idx]

    return {
        "text": most_relevant_text,
        "index": best_idx,
        "distance": distance
    }

def post(self, shared, prep_res, exec_res):
    shared["retrieved_document"] = exec_res

    print(f"📄 Retrieved document (index: {exec_res['index']}, distance: {exec_res['distance']})")
    print(f"📄 Most relevant text: \'{exec_res['text']}\'")

    return "default"

```

The `prep` gets query embedding and index, `exec` searches for the closest matching document chunk, and `post` saves the retrieved information.

## Node 6: The GenerateAnswerNode

This node creates the final answer using the retrieved context:

```

class GenerateAnswerNode(Node):
    def prep(self, shared):
        return shared["query"], shared["retrieved_document"]

    def exec(self, inputs):
        query, retrieved_doc = inputs

        prompt = f"""
Briefly answer the following question based on the context provided:
Question: {query}
Context: {retrieved_doc['text']}
Answer:

```

```

    answer = call_llm(prompt)

    return answer

def post(self, shared, prep_res, exec_res):
    shared["generated_answer"] = exec_res

    print("\n🤖 Generated Answer:")
    print(exec_res)

    return "default"

```

The `prep` gets the question and retrieved document, `exec` creates a prompt and calls an LLM, and `post` saves the generated answer.

## Connecting Everything Together

```

# Create offline flow for document processing
chunk_docs_node = ChunkDocumentsNode()
embed_docs_node = EmbedDocumentsNode()
create_index_node = CreateIndexNode()

# Connect nodes in sequence
chunk_docs_node >> embed_docs_node >> create_index_node
offline_flow = Flow(start=chunk_docs_node)

# Create online flow for answering questions
embed_query_node = EmbedQueryNode()
retrieve_doc_node = RetrieveDocumentNode()
generate_answer_node = GenerateAnswerNode()

# Connect nodes in sequence
embed_query_node >> retrieve_doc_node >> generate_answer_node
online_flow = Flow(start=embed_query_node)

```

This connects our nodes into two flows - one for document processing (run once) and one for answering questions (run for each query).

The complete working code for this tutorial is available at [GitHub: PocketFlow RAG Cookbook](#).

## Conclusion: The Beauty of RAG Simplicity

Now you understand the elegant simplicity of RAG systems:

1. **Offline Flow:** Process your documents once → Create vectors → Build searchable index
2. **Online Flow:** Process questions → Find relevant context → Generate grounded answers

This simple pattern powers even the most sophisticated RAG systems. Next time you encounter a complex RAG framework, look for these two workflows beneath the abstractions, and you'll understand how it works.

The power of RAG lies in its simplicity and effectiveness. By grounding AI responses in your specific documents, RAG dramatically improves accuracy, reduces hallucinations, and enables AI systems to work with your custom knowledge.

With what you've learned, you can now build your own RAG systems for any domain-specific knowledge base!

---

*Try Pocket Flow today and experience how 100 lines can replace hundreds of thousands! [GitHub Repository](#) | [Documentation](#) | [TypeScript Version](#)*

---

## Top comments (0)

[Code of Conduct](#) • [Report abuse](#)



Zachary Huang

Researcher @MSFTResearch AI Frontiers. I work on LLM Agents and Sys |  
Phd@ColumbiaCompSci | Prev: @GraySystemsLab @databricks | Fellowship: @GoogleAI

JOINED

Mar 4, 2025

### More from [Zachary Huang](#)

LLM Agents & Context: A Warrior's Guide to Navigating the Dungeon

#programming #python #beginners #tutorial

## LLM Agent's Arsenal: A Beginner's Guide to the Action Space

#programming #ai #llm #python

---

I just wrote a tutorial showing how to build an AI chatbot for your website that just works.

<https://dev.to/zachary62/the-easiest-way-to-build-an-ai-chatbot-for-your-website-full-dev-tutorial-37kp>

#ai #howto #tutorial #webdev

---